

EE345: Intro to Foundations of ML

M/W IEB 205 @ 10-11:20

Welcome



- Prof: Lillian Ratliff (ratliffl@uw)
 - Associate Professor ECE, Adjunct in CSE, AA;
 - Amazon Scholar
- Research: Machine learning, game theory, optimization

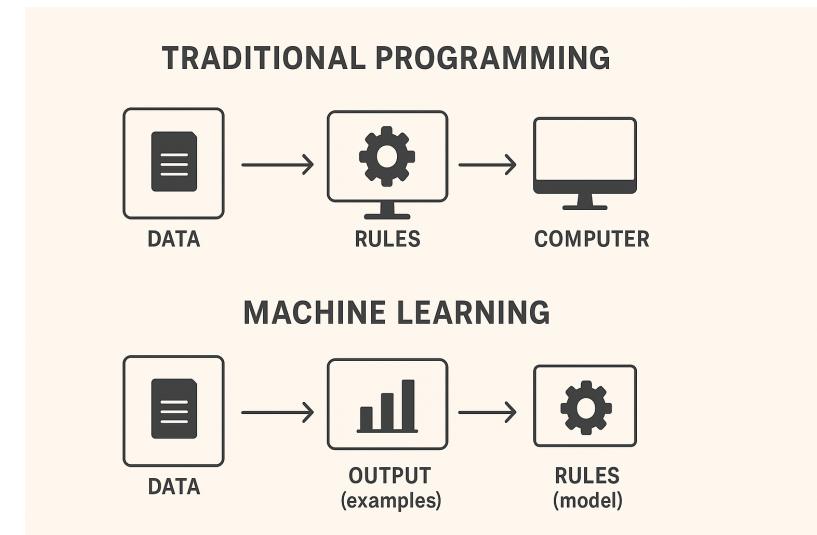


- TA: Adhyyan Narang (adhyyan@uw), ECE PhD Student
 - Research: Machine learning, game theory, optimization
-
- grader: Sergey Konyakin (serkon@uw), MS ECE Student

What is Machine Learning (ML)?

The process of giving machines the ability to *learn from data* so they can *recognize patterns* and make *predictions* or *decisions* without being explicitly programmed for every specific situation.

- The fundamental assumption of machine learning is that *pattern recognition is possible*. We take this as an axiom.
- The technical question: *What heuristics reasonably work to program computers to recognize those assumed patterns?*

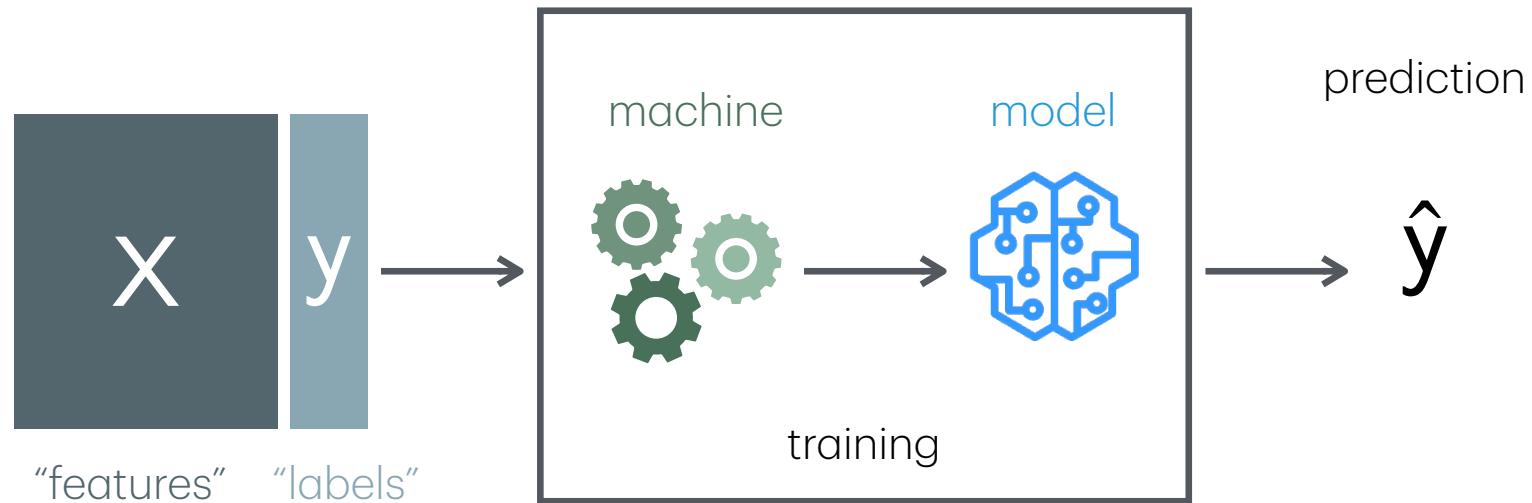


Ben Recht @ UC Berkeley

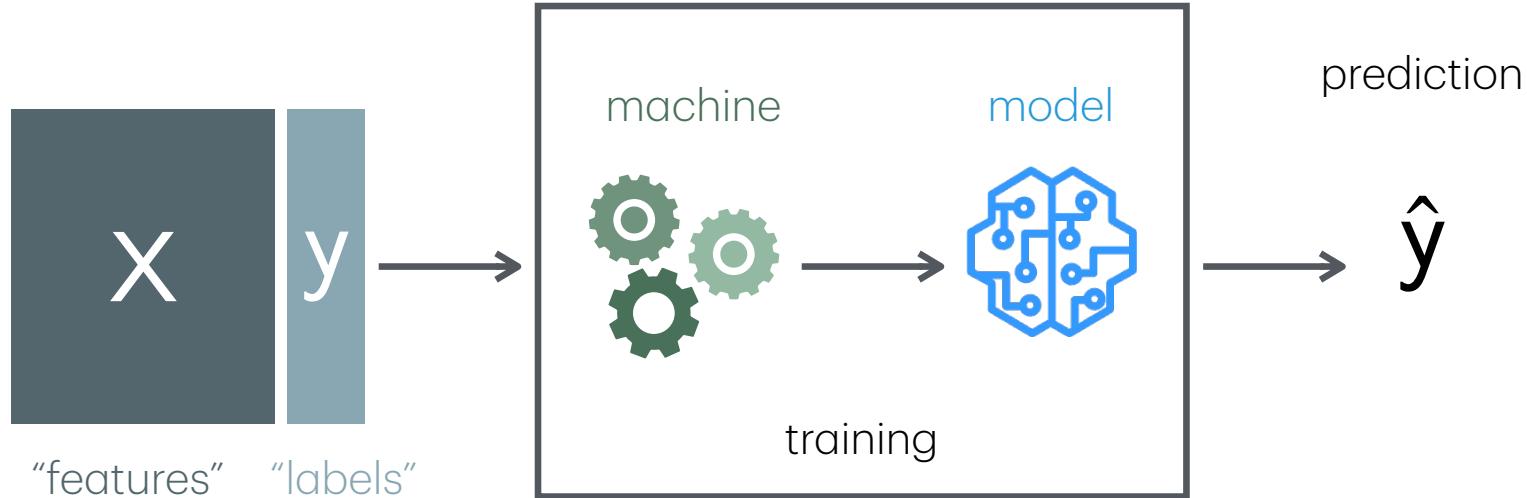
argmin
Machine Learning 101

What is Machine Learning (ML)?

The process of giving machines the ability to *learn from data* so they can *recognize patterns* and make *predictions* or *decisions* without being explicitly programmed for every specific situation.



ML examples



- house price \leftarrow [size, bedrooms, walkability score, ...]
- Blood pressure \leftarrow [age, weight, exercise, genetic factors, ...]
- exam score \leftarrow [hours studied, homework completed, sleep, ...]

$$\text{house price} = \theta_1 \cdot \text{size} + \theta_2 \cdot \text{bedrooms} + \theta_3 \cdot \text{walkability} + w$$

Course Focus

- Core/foundational topics:

- supervised learning (regression & classification),
- un-supervised (clustering, PCA)

- Advanced topics:

- kernel methods and non-linear feature construction
- generalization, regularization, cross-validation, bias-variance tradeoff
- loss function design/selection
- stochastic gradient descent,
- semi-supervised learning (PPI)
- feature analysis/design

- Advanced ML methods:

- Other ML Models: e.g., decision trees, neural networks, transformers, diffusion, reinforcement learning, multi-armed bandits
- Data/distribution shift (label vs concept vs covariate)

$$y = \theta_1 \cdot x_1 + \theta_2 \cdot x_2 + \dots + \theta_n \cdot x_n + w$$
$$y = \theta_1 \phi(x_1) + \dots + w$$
$$\log(x_1)$$

ML Foundations: Why?

This class will focus on foundations of ML such as linear and other simple, transparent models

ML Foundations: Why?

This class will focus on foundations of ML such as linear and other simple, transparent models

- Simple models → great for baselines, fast inference, small data, and transparency.
 - **building blocks:** Many complex architectures (deep nets, transformers) *reduce to stacks of linear predictors + nonlinearities.*
 - **give intuition:** understanding how coefficients scale features or how decision boundaries form makes it easier to grasp what deeper models are doing in high dimensions
 - **practical baselines:** in research and industry, you *always* start with a simple model to sanity-check.
 - **interpretable:** you can literally plot the data, draw the line or plane, and show predictions, which is much harder with a 70B-parameter LLM.

ML Foundations: Why?

This class will focus on foundations of ML such as linear and other simple, transparent models

- Simple models → great for baselines, fast inference, small data, and transparency.
 - **building blocks:** Many complex architectures (deep nets, transformers) *reduce to stacks of linear predictors + nonlinearities.*
 - **give intuition:** understanding how coefficients scale features or how decision boundaries form makes it easier to grasp what deeper models are doing in high dimensions
 - **practical baselines:** in research and industry, you *always* start with a simple model to sanity-check.
 - **interpretable:** you can literally plot the data, draw the line or plane, and show predictions, which is much harder with a 70B-parameter LLM.
- Complex models → great for scaling, raw data modalities, transfer learning, and long-term generalization.
 - **End-to-end:** less manual feature engineering
 - **Expressive:** can capture nonlinear, structured, and hierarchical patterns
 - **Reusable Architectures:** same backbone can be adapted across tasks
 - **Scales with Data/Compute:** often improves steadily as size grows

ML Foundations: Why?

This class will focus on foundations of ML such as linear and other simple, transparent models

ML Foundations: Why?

This class will focus on foundations of ML such as linear and other simple, transparent models

- Token Scoring (Softmax Head): **linear predictor** inside the LLM
 - At end of a transformer, there's a **linear layer** (weights + bias) that maps the hidden representation of a token into a large vector: $\text{logits} = W \cdot h_t + b$, $\Pr(\text{token} | h_t) = \text{softmax}(\text{logits})$

ML Foundations: Why?

This class will focus on foundations of ML such as linear and other simple, transparent models

- Token Scoring (Softmax Head): **linear predictor** inside the LLM
 - At end of a transformer, there's a **linear layer** (weights + bias) that maps the hidden representation of a token into a large vector: $\text{logits} = W \cdot h_t + b$, $\Pr(\text{token} | h_t) = \text{softmax}(\text{logits})$
- Classifier Head: LLM as a feature extractor —> **linear model** does the prediction.
 - In fine-tuning of GPT-like models (e.g., sentiment classification, toxicity detection, retrieval), often the LLM is frozen and a **simple linear predictor** “head” is trained on embeddings: e.g., Predictor = logistic regression that outputs “positive” or “negative” sentiment.

ML Foundations: Why?

This class will focus on foundations of ML such as linear and other simple, transparent models

- Token Scoring (Softmax Head): **linear predictor** inside the LLM
 - At end of a transformer, there's a **linear layer** (weights + bias) that maps the hidden representation of a token into a large vector: $\text{logits} = W \cdot h_t + b$, $\Pr(\text{token} | h_t) = \text{softmax}(\text{logits})$
- Classifier Head: LLM as a feature extractor —→ **linear model** does the prediction.
 - In fine-tuning of GPT-like models (e.g., sentiment classification, toxicity detection, retrieval), often the LLM is frozen and a **simple linear predictor** “head” is trained on embeddings: e.g., Predictor = logistic regression that outputs “positive” or “negative” sentiment.
- Reward Models in RLHF: reward assignment is just a **simple predictor**
 - preferences mapped to scalar reward via small head: **linear regression** (hidden state → scalar score).

ML Foundations: Why?

This class will focus on foundations of ML such as linear and other simple, transparent models

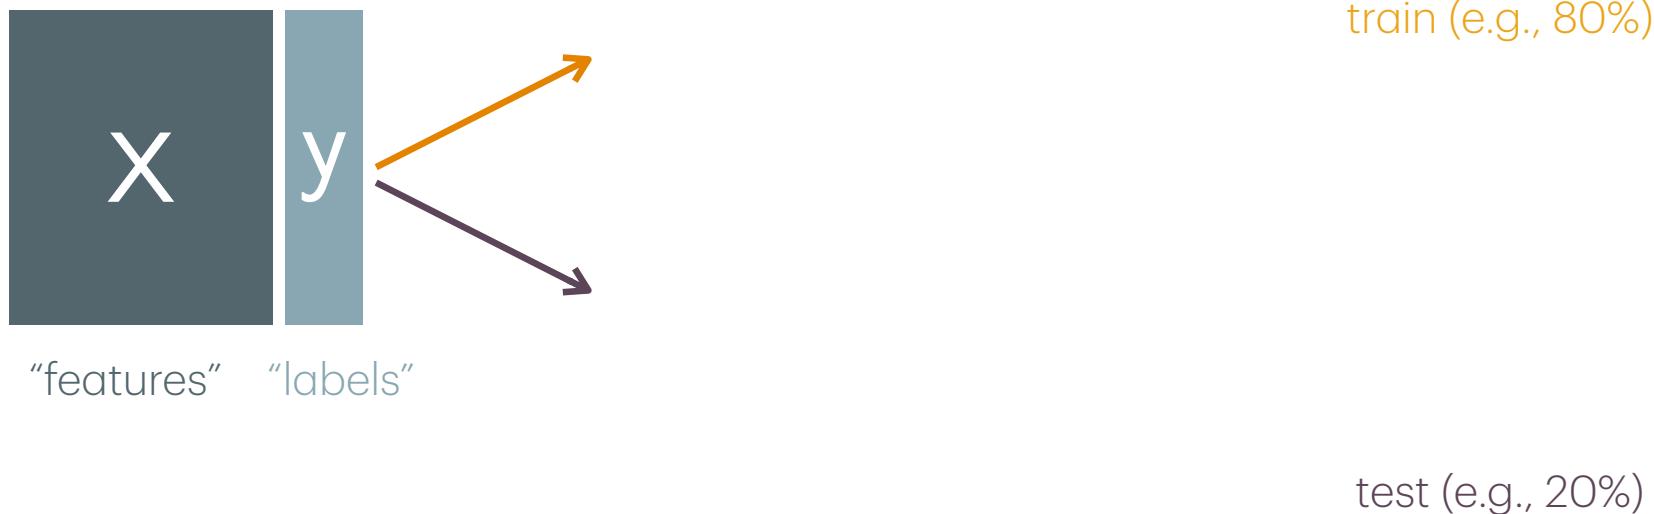
- Token Scoring (Softmax Head): **linear predictor** inside the LLM
 - At end of a transformer, there's a **linear layer** (weights + bias) that maps the hidden representation of a token into a large vector: $\text{logits} = W \cdot h_t + b$, $\Pr(\text{token} | h_t) = \text{softmax}(\text{logits})$
- Classifier Head: LLM as a feature extractor —→ **linear model** does the prediction.
 - In fine-tuning of GPT-like models (e.g., sentiment classification, toxicity detection, retrieval), often the LLM is frozen and a **simple linear predictor** “head” is trained on embeddings: e.g., Predictor = logistic regression that outputs “positive” or “negative” sentiment.
- Reward Models in RLHF: reward assignment is just a **simple predictor**
 - preferences mapped to scalar reward via small head: **linear regression** (hidden state → scalar score).
- Routing in LLMs: gating via **linear or softmax classifiers**
 - LLMs (like Switch Transformers, Mixture-of-Experts) use a **gate network** to decide which expert block to send a token through

ML Foundations: Why?

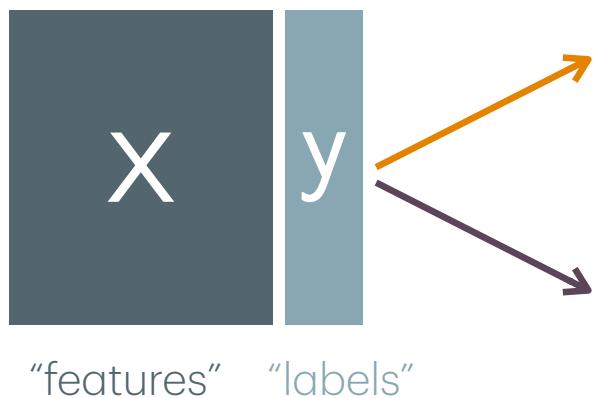
This class will focus on foundations of ML such as linear and other simple, transparent models

- Token Scoring (Softmax Head): **linear predictor** inside the LLM
 - At end of a transformer, there's a **linear layer** (weights + bias) that maps the hidden representation of a token into a large vector: $\text{logits} = W \cdot h_t + b$, $\Pr(\text{token} | h_t) = \text{softmax}(\text{logits})$
- Classifier Head: LLM as a feature extractor —→ **linear model** does the prediction.
 - In fine-tuning of GPT-like models (e.g., sentiment classification, toxicity detection, retrieval), often the LLM is frozen and a **simple linear predictor** “head” is trained on embeddings: e.g., Predictor = logistic regression that outputs “positive” or “negative” sentiment.
- Reward Models in RLHF: reward assignment is just a **simple predictor**
 - preferences mapped to scalar reward via small head: **linear regression** (hidden state → scalar score).
- Routing in LLMs: gating via **linear or softmax classifiers**
 - LLMs (like Switch Transformers, Mixture-of-Experts) use a **gate network** to decide which expert block to send a token through
- Internal probing via linear classifiers enables decoding (e.g., word knowledge, hidden features)
 - internally, the model is arranging information so that simple predictors can pull it out

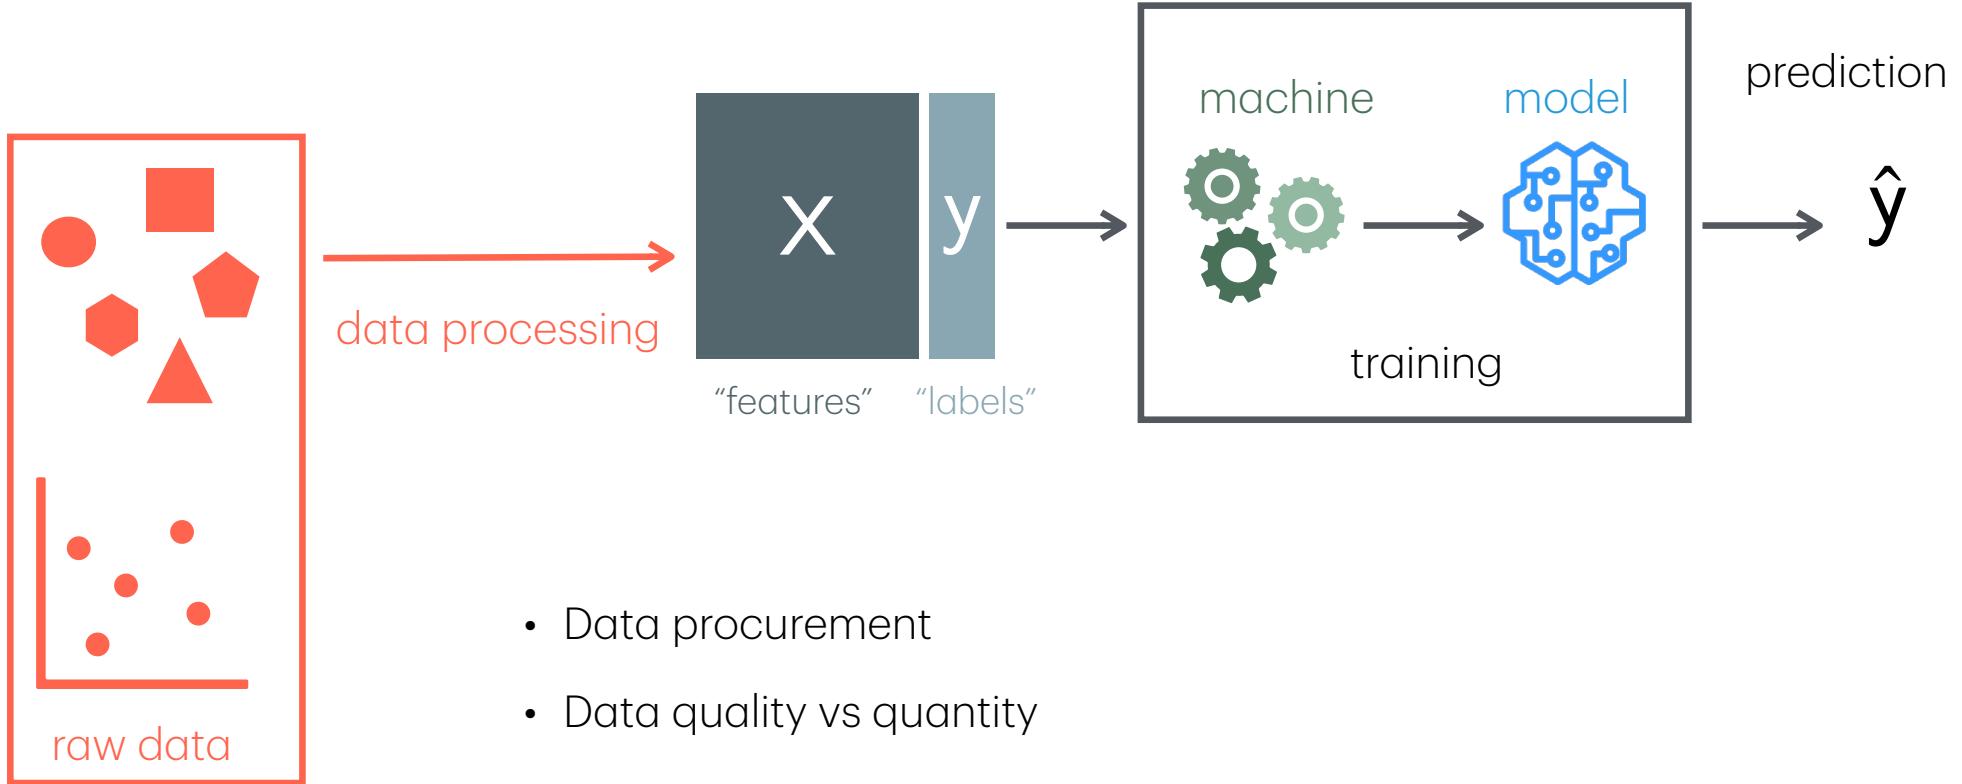
How to evaluate ML?: Hold-out method



How to evaluate ML?: Hold-out method

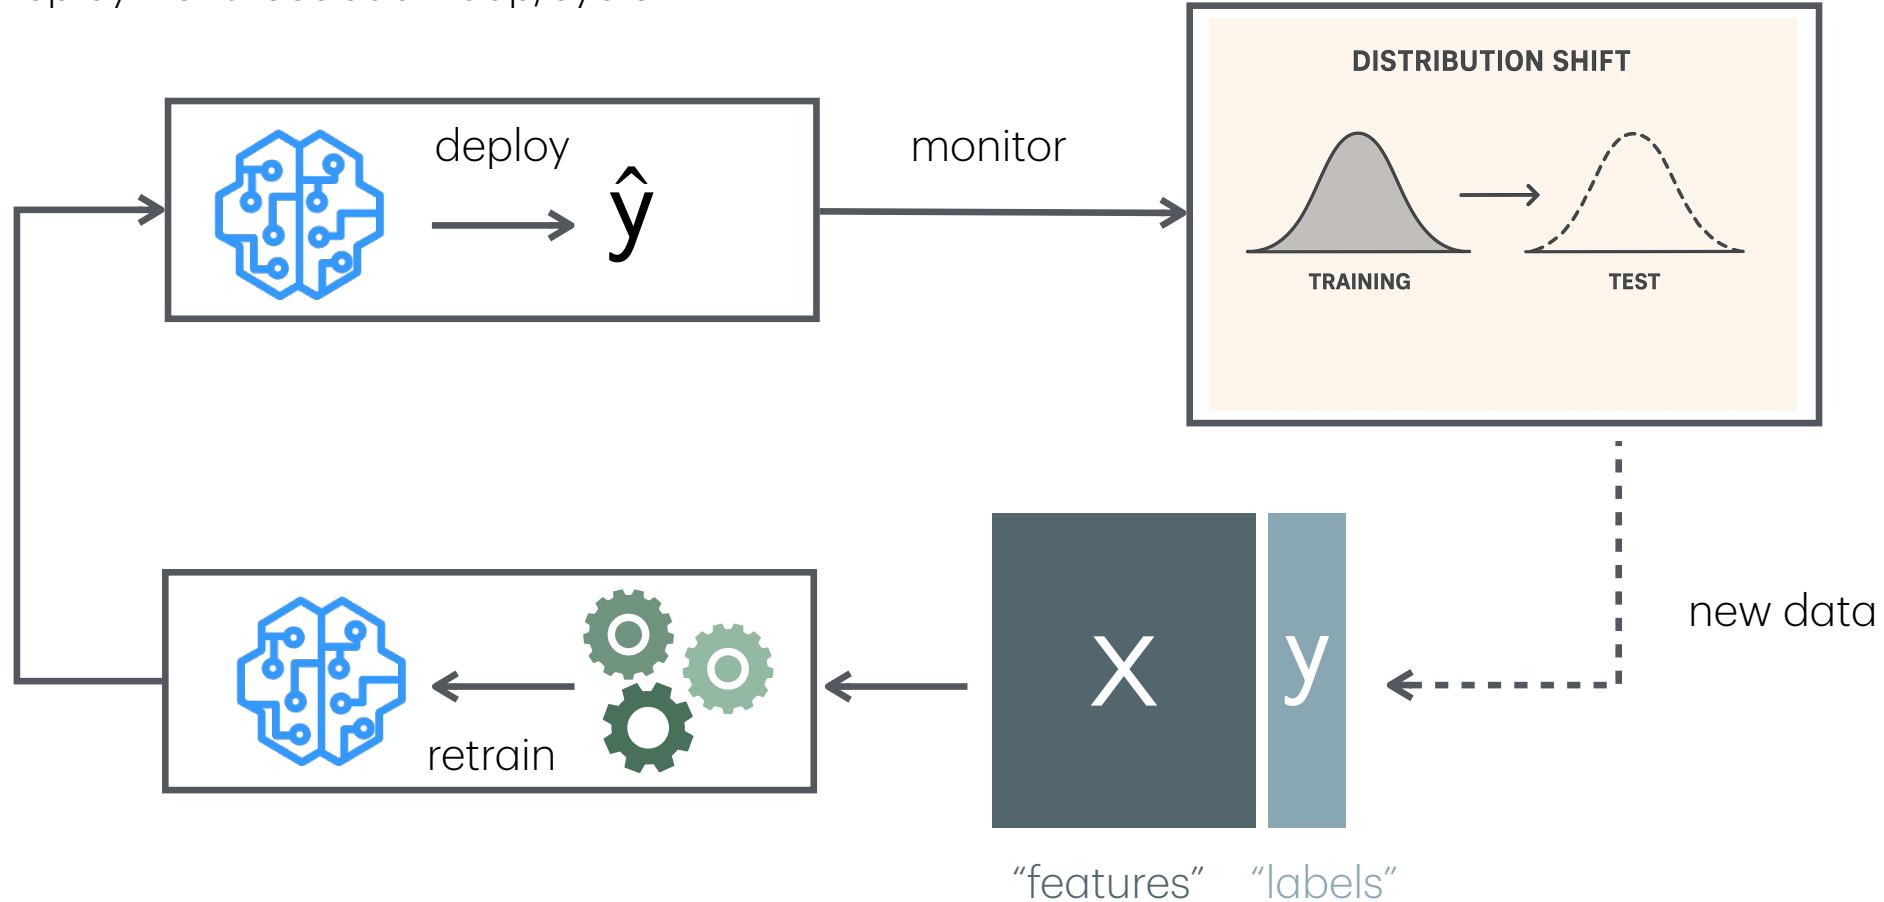


What is often overlooked in ML courses?

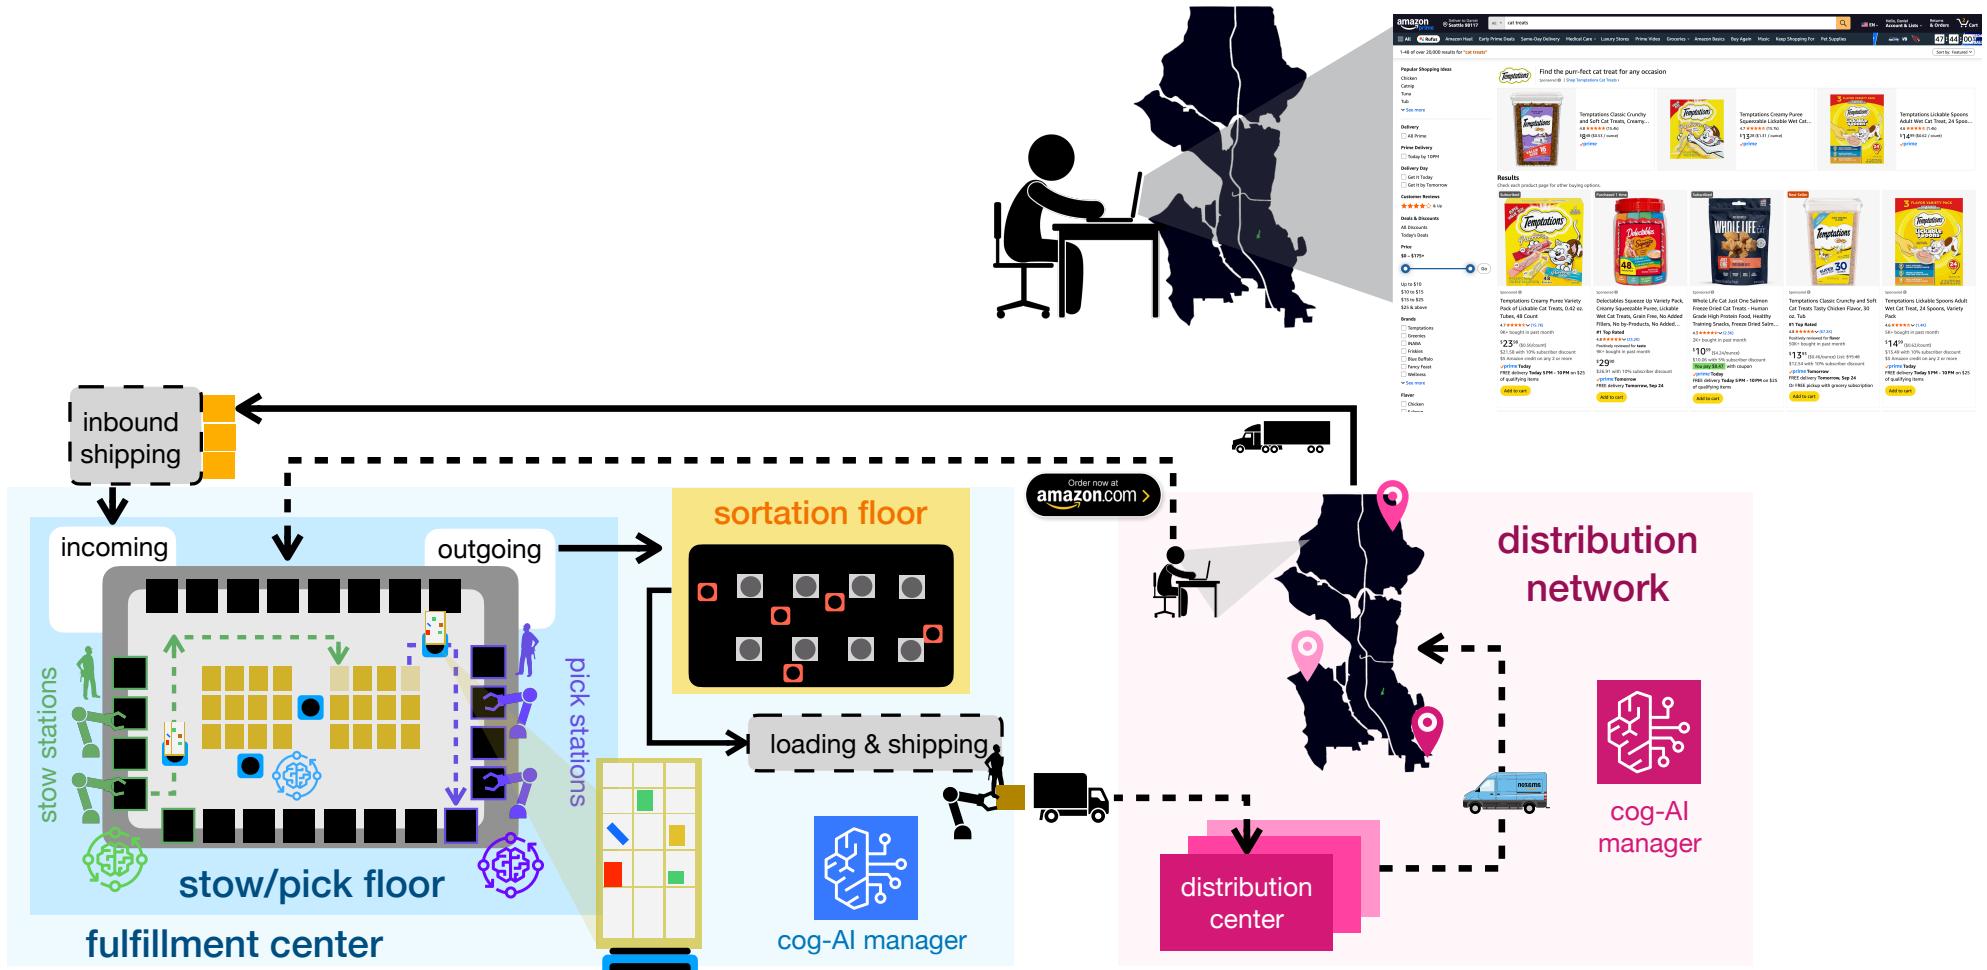


What is often overlooked in ML courses?

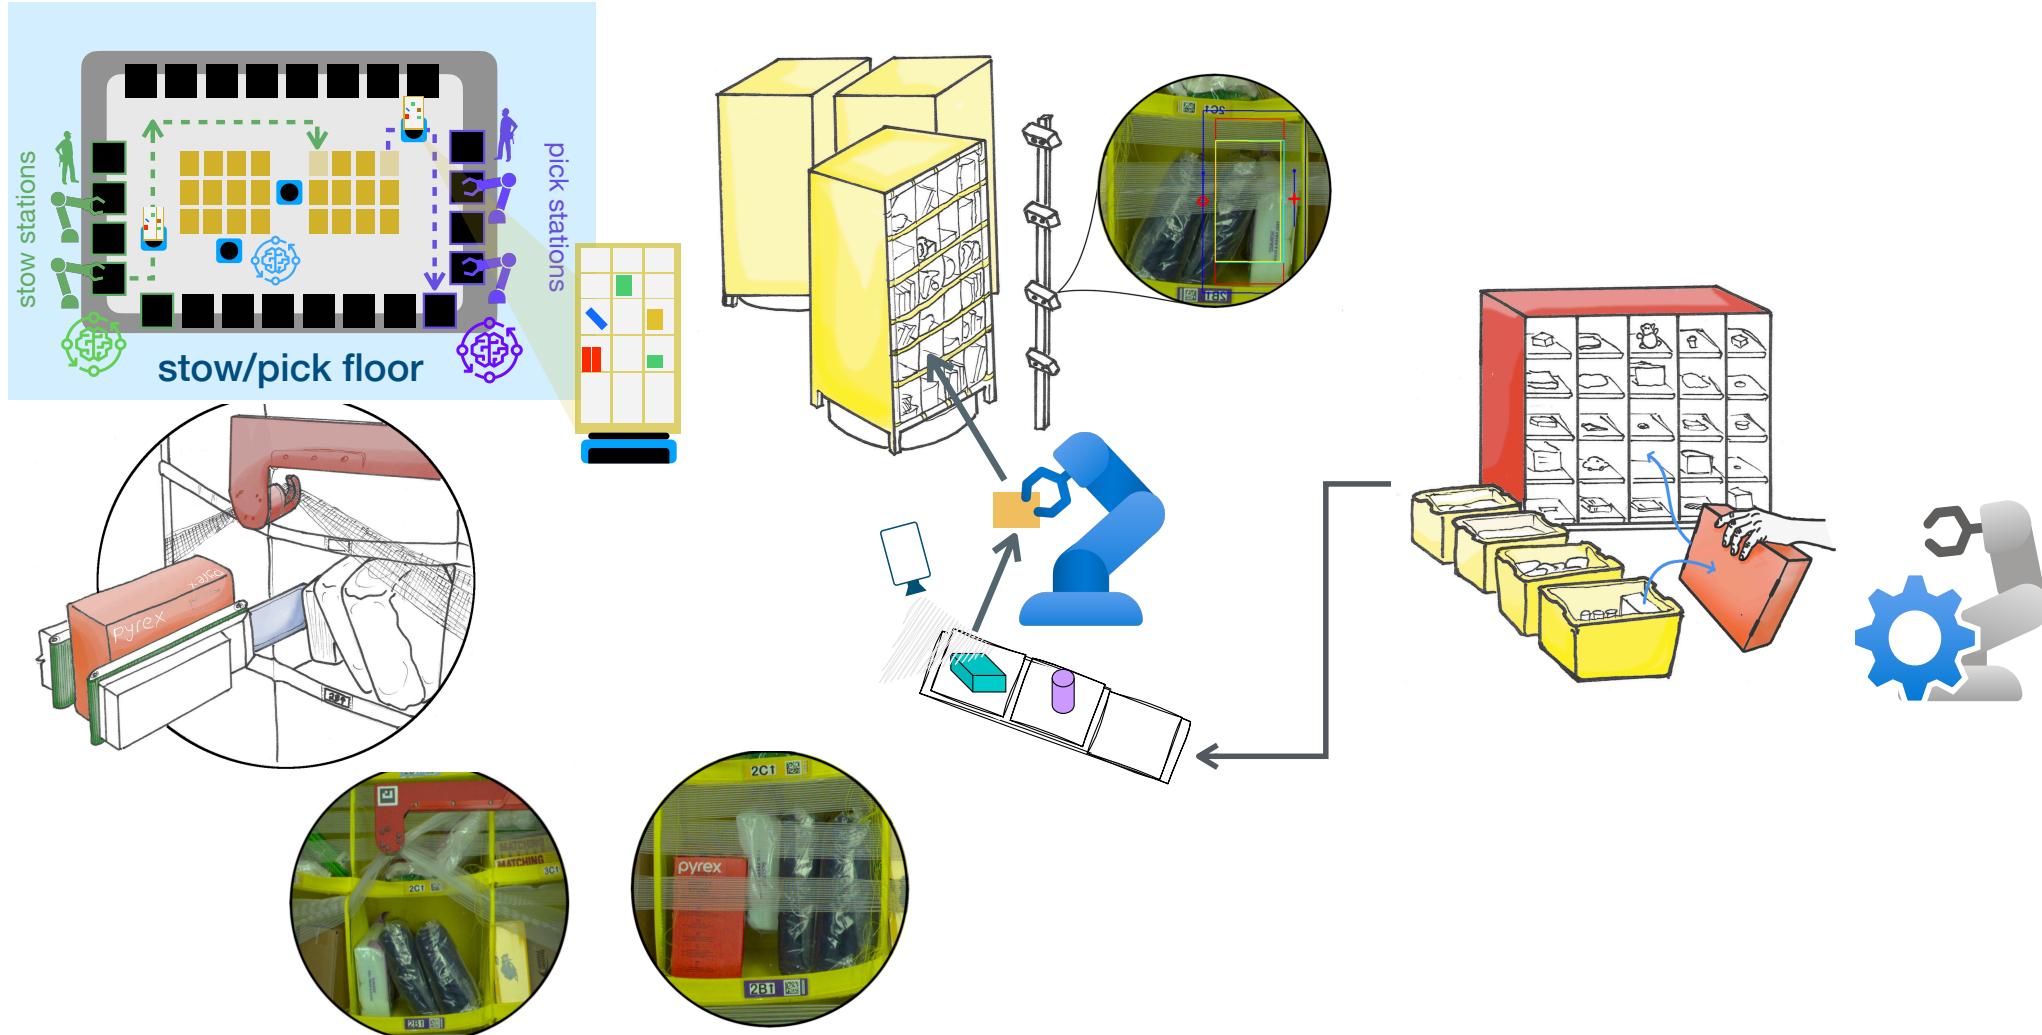
- Deployment feedback loop/cycle



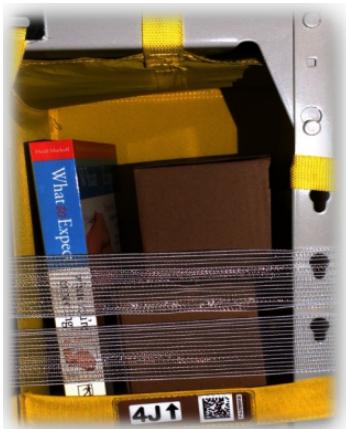
Motivation: Fulfillment in E-Commerce



Motivation: Robotic Stowing (Embodied AI)



Motivation: Robotic Stowing (Embodied AI)



Motivation: Robotic Stowing (Embodied AI)



Motivation: Robotic Stowing (Embodied AI)



Motivation: Robotic Stowing (Embodied AI)



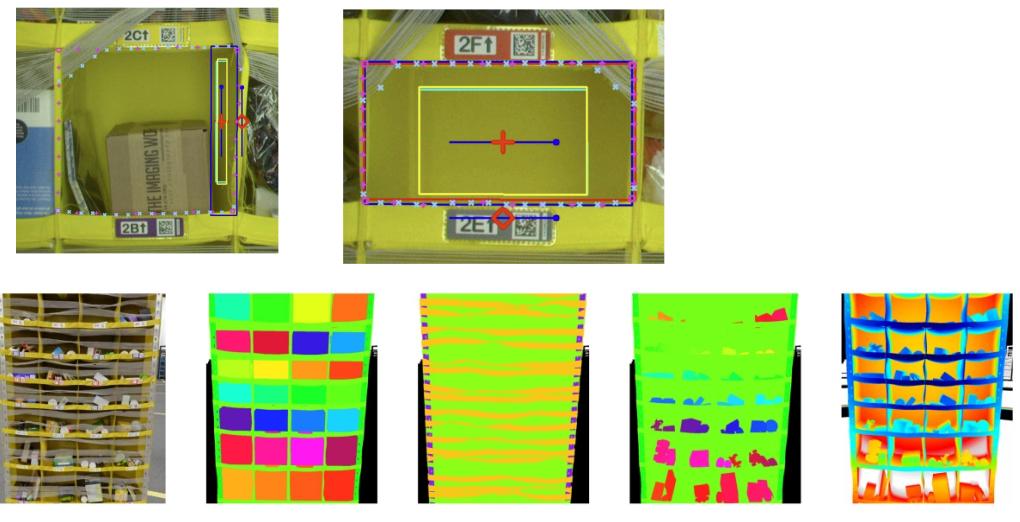
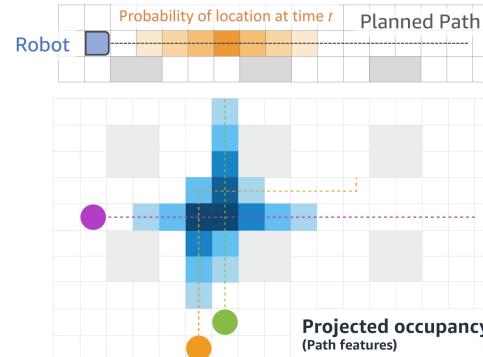
Motivation: Robotic Stowing (Embodied AI)



Motivation: What are the ML problems?

Motivation: What are the ML problems?

- Forecasting congestion (regression)
- Space/sweep length prediction (regression)
- Predicting success/outcomes (classification)
- Amnesty/defects (classification)
- Synthetic data (generative models)
- Matching policies (RL/Q learning)
- Path planning (reinforcement learning)



ML Ops (Machine learning operations)

- Data quality vs quantity
- Data imbalance
- Performance monitoring & retraining
- Distribution Shift
- Latency

Course Focus

- Core/foundational topics:
 - supervised learning (regression & classification),
 - un-supervised (clustering, PCA)
- Advanced topics:
 - kernel methods and non-linear feature construction
 - generalization, regularization, cross-validation
 - stochastic gradient descent,
 - semi-supervised learning (PPI)
 - feature analysis/design
- Advanced ML methods:
 - Other ML Models: e.g., decision trees, neural networks, transformers, diffusion
 - Data/distribution shift (label vs concept vs covariate)

Course Information

- Website: ratliffj.github.io/F25-EE345/
- Canvas: canvas.uw.edu/courses/1860120
 - quizzes used for self-grading hw
 - submitting hw and Python exercises
- Chat (Discord?/Canvas?)
- Gradescope: used for exam grading
- Zoom: washington.zoom.us/j/95494709719
 - If you need to miss class for some reason, there will also be a zoom room to join remotely

Course Structure/Logistics

- Lectures (M/W IEB 205 @ 10-11:20)
- Discussions (F IEB 205 @ 9:30-11:20)
- Communication/Questions
 - Online Chat (Discord?/Canvas?)
 - Office Hours: TBD
- Guest Speakers from Industry!
- Grade distribution: TBD
- Books?
 - Boyd & Vandenberghe: Intro to Applied LA
 - Calafiore & El Ghaoui: Optimization Models
 - James, Witten, Hastie, Tibshirani: An intro to stat learning
- Homework:
 - weekly (~6-7 total);
 - released W; due following W
 - self-graded & due following M (grader grades 1 problem @ random)
- Python Exercises:
 - weekly assignments (~6-7 total), discussion section will be used to do majority of the work
 - group project: end of term
- Exams:
 - Midterm (~5-6th week)
 - Final (finals week)

Comments on AI Tool Usage

- **Goal:** Teach everyone best practice on how to use AI tools for coding efficiently.
- **Question everything!** They are wrong often, hallucinate frequently, and sycophantic to a fault!
 - If you get an answer you don't understand or trust, don't be afraid to ask us

Cursor AI

Up Next...

- Next week we will start with
 - data fitting and simple regression problems
 - data structure
 - basic linear algebra constructs (matrices, vectors, norms) for ML modelling